

SHARED LIBRARY OPERATORS IN DORA

Creating, linking, calling into shared libraries

Philipp Oppermann

os2edu

2023-11-17

Agenda

- Background: Executables and libraries
- Motivation: Operators in Dora
- Shared Libraries
 - C and Rust examples
 - Shared library search paths
 - Dynamic loading
- Improving Safety

Executables and Libraries

- Executables are programs with a main function
- Libraries provide callable functions to executable and other libraries
 - e.g. Rust crates can use libraries through a `[dependencies]` section in their `Cargo.toml` file
 - typical file endings of libraries are `.a`, `.lib`, `.so`, `.dll`

Executables and Libraries

- Executables are programs with a main function
- Libraries provide callable functions to executable and other libraries
 - e.g. Rust crates can use libraries through a [dependencies] section in their Cargo.toml file
 - typical file endings of libraries are .a, .lib, .so, .dll

Example: Use readelf to read type of ELF files:

```
> readelf --file-header /usr/bin/ls | grep "Type"  
Type:                                DYN (Position-Independent Executable file)  
> readelf --file-header /usr/lib32/libc.so.6 | grep "Type"  
Type:                                DYN (Shared object file)
```

C Example

- executable.c:

```
#include <stdio.h>
void main() {
    printf("hello world");
}
```

- create executable: **gcc -o executable executable.c**
- run executable: **./executable** → prints "hello world"

C Example

- executable.c:

```
#include <stdio.h>
void main() {
    printf("hello world");
}
```

- create executable: **gcc -o executable executable.c**
- run executable: **./executable** → prints "hello world"

- sum.c:

```
int sum(const int a, const int b) {
    return a + b;
}
```

- try to build as executable using **gcc -o sum sum.c** → error: undefined reference to `main'
- create object file: **gcc -c sum.c**
- then create static library: **ar r libsum.a sum.o**

C Example: Using static library

- Use libsum.a in executable.c:
 - define header file with function signature in sum.h:

```
int sum(const int a, const int b);
```

- include sum.h in executable:

```
#include <stdio.h>
#include "sum.h"

void main() {
    printf("hello world %i", sum(40, 2));
}
```

- create executable, linking libsum.a: **gcc -o executable executable.c -L. -lsum**
- run executable: **./executable** → prints "hello world 42"

C Example: Using static library

- Use libsum.a in executable.c:
 - define header file with function signature in sum.h:

```
int sum(const int a, const int b);
```

- include sum.h in executable:

```
#include <stdio.h>
#include "sum.h"

void main() {
    printf("hello world %i", sum(40, 2));
}
```

- create executable, linking libsum.a: `gcc -o executable executable.c -L. -lsum`
 - run executable: `./executable` → prints "hello world 42"
- **Note:** The argument order is important: `-l` needs to come after `executable.c`
 - else **undefined reference to 'sum'** error, as linker only looks for symbols that are referenced

Rust Example

- Executable using cargo:

```
> cargo new --bin executable
  Created binary (application) 'executable' package
> cd executable
> cargo build
  Compiling executable v0.1.0 (/../executable)
  Finished dev [unoptimized + debuginfo] target(s) in 0.13s
> target/debug/executable
Hello, world!
```

Rust Example

- Executable using cargo:

```
> cargo new --bin executable
Created binary (application) 'executable' package
> cd executable
> cargo build
Compiling executable v0.1.0 (/../executable)
Finished dev [unoptimized + debuginfo] target(s) in 0.13s
> target/debug/executable
Hello, world!
```

- Library using cargo:
 - create: `cargo new --lib sum`
 - modify `src/lib.rs`:

```
pub fn sum(left: isize, right: isize) → isize {
    left + right
}
```

- build using `cargo build` → creates `target/debug/libsum.rlib` library

Rust Example: Using Rust library

- Using libsum.rlib in executable crate:
 - add dependency in executable/Cargo.toml:

```
[dependencies]
sum = { path = "../sum" }
```

- use sum function in executable/src/main.rs:

```
use sum::sum;

fn main() {
    println!("Hello, world! {}", sum(40, 2));
}
```

- rebuild using cargo build
- run target/debug/executable → outputs "Hello, world! 42"

Note: [rlib libraries](#) are Rust-specific and might change over time

Operators in Dora

- Dora operators are components that can be executed by the Dora runtime
 - act as nodes in the dataflow graph → receive inputs and send outputs
 - a single Dora runtime process can run multiple operators concurrently
- Operators are **libraries** that implement a specific template
 - e.g. they need to provide an **on_event** function

Operators in Dora

- Dora operators are components that can be executed by the Dora runtime
 - act as nodes in the dataflow graph → receive inputs and send outputs
 - a single Dora runtime process can run multiple operators concurrently
- Operators are **libraries** that implement a specific template
 - e.g. they need to provide an **on_event** function

Challenge:

- We don't want to recompile the Dora runtime when loading operators
- `.a` and `.rlib` libraries need to be included at compile time

→ we need to use a different library format that can be loaded without recompiling

Python Libraries

- Python is an interpreted language → no precompiled code
- Rust-based Dora runtime can import Python operators using pyo3 crate, e.g.:

```
Python::with_gil(|py| {  
    let sum = PyModule::from_code(py, "  
        def sum(x, y):  
            return x + y  
    ", "sum.py", "sum")?;  
  
    let result: i32 = sum.getattr("sum")?.call2((40, 2))?.extract()?;  
    assert_eq!(result, 42);  
  
    Ok::<(), PyErr>(())  
})
```

- pyo3 returns errors when operator has wrong format, e.g. is missing a required function

→ we want something similar for compiled languages: **shared libraries**

Shared libraries

- Approach
 - Precompile library as before, but include additional metadata
 - Compile executable against stub library that describes the template
 - When executable is loaded, combine compiled executable and library using the metadata
- Supported and used on all major OS platforms, but different format:
 - Linux: `.so`
 - Windows: `.dll`
 - MacOS: `.dylib`

Shared libraries

- Approach
 - Precompile library as before, but include additional metadata
 - Compile executable against stub library that describes the template
 - When executable is loaded, combine compiled executable and library using the metadata
- Supported and used on all major OS platforms, but different format:
 - Linux: `.so`
 - Windows: `.dll`
 - MacOS: `.dylib`
- Our example Rust executable already has some shared library dependencies:

```
> readelf --dynamic target/debug/executable | grep "Shared"
0x0000000000000001 (NEEDED)           Shared library: [libgcc_s.so.1]
0x0000000000000001 (NEEDED)           Shared library: [libc.so.6]
0x0000000000000001 (NEEDED)           Shared library: [ld-linux-x86-64.so.2]
```


Example: C Shared library

- We reuse the same `executable.c`, `sum.h`, and `sum.c` as before

```
int sum(const int a, const int b) {  
    return a + b;  
}
```

- Build steps:
 - create position-independent object file for sum: `gcc -c -fPIC sum.c`
 - then create shared library: `gcc -shared sum.o -o shared/libsum.so`
 - link against the shared library: `gcc -o executable executable.c -Lshared -lsum`

Example: C Shared library

- We reuse the same `executable.c`, `sum.h`, and `sum.c` as before

```
int sum(const int a, const int b) {  
    return a + b;  
}
```

- Build steps:
 - create position-independent object file for sum: `gcc -c -fPIC sum.c`
 - then create shared library: `gcc -shared sum.o -o shared/libsum.so`
 - link against the shared library: `gcc -o executable executable.c -Lshared -lsum`
- Try running it:
 - > `./executable`
 - `./executable: error while loading shared libraries: libsum.so: cannot open shared object file: No such file or directory`
 - shared library is required for running the executable, but not found

Shared Library Search Paths

Lookup of shared libraries depends on operating system:

- On Linux, the linker tries the following directories:
 1. directories listed in the LD_LIBRARY_PATH environment variable
 2. paths specified in the executable itself (through an [rpath](#) attribute)
 3. system search paths
- MacOS is similar, but it uses an env variable called DYLD_LIBRARY_PATH
- On Windows, the behavior is [more complex](#). Some differences are:
 - no separate environment variable, instead the normal PATH is searched
 - the folder containing the executable is searched
 - the *current working directory* is searched too by default (note: this can be dangerous)

Fixing the C Shared Library Example

- Problem: The `libsum.so` library lives in the shared subdirectory
 - not in the default search path

Fixing the C Shared Library Example

- Problem: The `libsum.so` library lives in the shared subdirectory
 - not in the default search path
- Solution 1: Set `LD_LIBRARY_PATH` environment variable:

```
> LD_LIBRARY_PATH=shared ./executable  
hello world 42
```

Fixing the C Shared Library Example

- Problem: The `libsum.so` library lives in the shared subdirectory
 - not in the default search path
- Solution 1: Set `LD_LIBRARY_PATH` environment variable:

```
> LD_LIBRARY_PATH=shared ./executable  
hello world 42
```

- Solution 2: Move library to system search path → not recommended

Fixing the C Shared Library Example

- Problem: The `libsum.so` library lives in the shared subdirectory
 - not in the default search path
- Solution 1: Set `LD_LIBRARY_PATH` environment variable:

```
> LD_LIBRARY_PATH=shared ./executable  
hello world 42
```

- Solution 2: Move library to system search path → not recommended
- Solution 3: Set `rpath` attribute when building executable:

```
> gcc -o executable executable.c -Lshared -lsum -Wl,-rpath shared  
> ./executable  
hello world 42
```

Example: Rust Shared Library

- Steps:

- `cargo new --lib sum2`
- modify `sum2/src/lib.rs`:

```
pub extern "C" fn sum(left: isize, right: isize) → isize {  
    left + right  
}
```

- set the crate-type in `Cargo.toml`:

```
[lib]  
crate-type = ["cdylib"]
```

- run `cargo build` to create `libsum2.so` in `target/debug`

Example: Rust Shared Library

- Steps:

- `cargo new --lib sum2`
- modify `sum2/src/lib.rs`:

```
pub extern "C" fn sum(left: isize, right: isize) → isize {  
    left + right  
}
```

- set the crate-type in `Cargo.toml`:

```
[lib]  
crate-type = ["cdylib"]
```

- run `cargo build` to create `libsum2.so` in `target/debug`
- Created shared library *should be* C-compatible, but linking with executable somehow fails:

```
> gcc -o executable executable.c -Lsum2/target/debug -lsum2  
/usr/bin/ld: /tmp/ccf95YLL.o: in function `main':  
executable.c:(.text+0x13): undefined reference to `sum'
```

Debug Rust Shared Library Example

- Print available symbols using nm:

```
> nm -gD sum/target/debug/libsum.so  
w __cxa_finalize@GLIBC_2.2.5  
w __gmon_start__  
w _ITM_deregisterTMCloneTable  
w _ITM_registerTMCloneTable
```

→ no sum function

Debug Rust Shared Library Example

- Print available symbols using nm:

```
> nm -gD sum/target/debug/libsum.so  
w __cxa_finalize@GLIBC_2.2.5  
w __gmon_start__  
w _ITM_deregisterTMCloneTable  
w _ITM_registerTMCloneTable
```

→ no sum function

- Reasons:
 - Rust functions are not exported as symbols by default
 - Rust function names are *mangled* by default to prevent name conflicts

Debug Rust Shared Library Example

- Print available symbols using nm:

```
> nm -gD sum/target/debug/libsum.so  
w __cxa_finalize@GLIBC_2.2.5  
w __gmon_start__  
w _ITM_deregisterTMCloneTable  
w _ITM_registerTMCloneTable
```

→ no sum function

- Reasons:
 - Rust functions are not exported as symbols by default
 - Rust function names are *mangled* by default to prevent name conflicts
- Solution: Set `#[no_mangle]` attribute:

```
#[no_mangle]  
pub extern "C" fn sum(left: isize, right: isize) → isize { ... }
```

→ after a cargo build, the gcc link error is now fixed

Link Shared Library from Rust

- Rust does not support header files
 - instead, use extern block to specify dependency on external sum function:

```
// in executable/src/main.rs
extern "C" {
    fn sum(a: isize, b: isize) -> isize;
}
```

- the [bindgen](#) crate allows auto-generating this extern block based on header files

Link Shared Library from Rust

- Rust does not support header files
 - instead, use extern block to specify dependency on external sum function:

```
// in executable/src/main.rs
extern "C" {
    fn sum(a: isize, b: isize) -> isize;
}
```

- the [bindgen](#) crate allows auto-generating this extern block based on header files
- Calling external functions is unsafe, as the Rust compiler cannot guarantee their safety:

```
println!("Hello, world! {}", unsafe { sum(40, 2) });
```

Link Shared Library from Rust

- Rust does not support header files
 - instead, use extern block to specify dependency on external sum function:

```
// in executable/src/main.rs
extern "C" {
    fn sum(a: isize, b: isize) → isize;
}
```

- the [bindgen](#) crate allows auto-generating this extern block based on header files
- Calling external functions is unsafe, as the Rust compiler cannot guarantee their safety:

```
println!("Hello, world! {}", unsafe { sum(40, 2) });
```

- Pass linker flags via new build.rs build script:

```
fn main() {
    println!("cargo:rustc-link-search=native=../sum2/target/debug");
    println!("cargo:rustc-link-lib=dylib=sum2");
}
```

→ cargo build now links to the shared library

Shared libraries: Pros and Cons

Advantages:

- Avoid duplication of common libraries
 - Both on disk and after loading in RAM
 - Example: `libLLVM-15.so` is over 100MiB
- Security: Updating a shared library fixes all executables
 - Example: SSL libraries
- Platform-specific libraries
 - Example: `libc`

Drawbacks:

- Program might not work when library is missing or at wrong version
- Distribution is more difficult

Shared Library Operators in Dora

Challenges:

- The dora runtime is compiled ahead of time
 - we don't have access to the operator at this point → workaround possible using stub library
- We want to load multiple operators → name conflicts
 - e.g., there can only be one `on_event` function
- The dora runtime should be able to recover from load errors
 - operators in an incompatible format should not bring the runtime down
 - load errors should result in useful error messages

→ use dynamic loading to link library manually at runtime

Dynamic Loading

- Don't link to the shared library directly
- Instead, load library at runtime using system functions
 - on Linux and MacOS: `dlopen`, `dlsym`, `dlclose`
 - on Windows: `LoadLibraryExW`, `GetProcAddress`, `FreeLibrary`

Dynamic Loading

- Don't link to the shared library directly
- Instead, load library at runtime using system functions
 - on Linux and MacOS: `dlopen`, `dlsym`, `dlclose`
 - on Windows: `LoadLibraryExW`, `GetProcAddress`, `FreeLibrary`
- Advantages:
 - more resilient → better error messages
 - no name conflicts → multiple operators can be loaded simultaneously
 - no stubs needed for building Dora runtime
- Drawbacks: less convenient, more manual work necessary

Example: Dynamic Loading using dlopen

```
#include <stdio.h>
#include <dlfcn.h>

int main(void) {
    void *handle = dlopen("libsum.so", RTLD_LAZY);
    if (!handle) { return -1; }

    int (*sum)(int, int) = (int (*)(int, int)) dlsym(handle, "sum");
    if (dlerror() != NULL) { return -1; }

    printf("%d\n", (*sum)(40, 2));

    dlclose(handle);
    return 0;
}
```

Dynamic Loading in Rust

Use the cross-platform libloading crate:

```
let result = unsafe {  
    let lib = libloading::Library::new("shared/libsum.so");  
    let func: libloading::Symbol<unsafe extern fn(usize, usize) -> isize> = lib.get(b"sum");  
    func(40, 2)  
};  
assert_eq!(result, 42);
```

Dynamic Loading in Rust

Use the cross-platform libloading crate:

```
let result = unsafe {  
    let lib = libloading::Library::new("shared/libsum.so"?;  
    let func: libloading::Symbol<unsafe extern fn(isize, isize) → isize> = lib.get(b"sum"?;  
    func(40, 2)  
};  
assert_eq!(result, 42);
```

Safety:

- We get nice error messages when the library or the requested symbol don't exist
 - e.g. "Error: shared/libsum.so: cannot open shared object file: No such file or directory"
 - e.g. "Error: shared/libsum.so: undefined symbol: sum"
- But there is no type checking → the function signature is not verified
 - that's the reason why all of the above code is considered unsafe

Improving Safety

- Provide header files to check operator signature
 - Still very easy to cause undefined behavior accidentally

Improving Safety

- Provide header files to check operator signature
 - Still very easy to cause undefined behavior accidentally
- Use the [abi_stable](#) crate
 - designed for Rust-to-Rust FFI
 - generates static variables with type and layout information → verify on load
 - provides FFI-compatible wrappers for standard library types
 - custom types can [derive the StableAbi trait](#)
 - even provides trait object support through [DynTrait](#) (for a selection of traits)

Improving Safety

- Provide header files to check operator signature
 - Still very easy to cause undefined behavior accidentally
- Use the [abi_stable](#) crate
 - designed for Rust-to-Rust FFI
 - generates static variables with type and layout information → verify on load
 - provides FFI-compatible wrappers for standard library types
 - custom types can [derive the StableAbi trait](#)
 - even provides trait object support through [DynTrait](#) (for a selection of traits)
- `abi_stable` guards against accidental incompatibilities, but not against malicious inputs
 - → isolate untrusted shared library operators into separate runtime process

Future Work: WebAssembly Operators

- WebAssembly is a sandboxed executable format designed for the web
 - even safer than Python (no dangerous FFI calls possible)
- Many languages can be compiled to WebAssembly, including Rust, C, etc.
- There are multiple mature Rust-based WebAssembly runtimes available
 - [Wasmtime](#) based on Cranelift
 - [WasmEdge](#) based on LLVM

→ we plan to add support for WebAssembly operators to Dora in the future

Summary

- Creating shared libraries in C and Rust
 - Different shared library search algorithms on Linux and Windows
- Advantages and drawbacks of shared libraries
 - avoid duplication, faster security fixes, use platform-specific code
 - missing or incompatible libraries prevent running of executable
- Dynamic loading
 - using `dlopen/LoadLibraryExW` or the `libloading` crate
 - enables independent linking, better error messages, and avoids name conflicts
 - drawback: manual type specifications needed → mistakes can cause undefined behavior
- Improving safety using `abi_stable` crate → future work: sandboxed WebAssembly operators