

# **CUSTOM SERIALIZATION WITH SERDE**

**Supporting the ROS2 message format in Dora**

Philipp Oppermann

os2edu

2023-11-24

# Agenda

- Introduction to serde crate
  - Simple serialization using derive attributes
  - The Serialize and Deserialize traits in detail
  - Example implementations
- Reading and writing ROS2 messages in Dora
  - Serializer with dynamic target type
  - Dynamic deserialization using DeserializeSeed

# The serde crate

- Framework for serializing and deserializing Rust data structures
- Based on Rust's trait system
- Independent of data format
  - Generic data format that can represent all Rust types
  - External crates to translate serde data format into target format
    - E.g. for JSON, YAML, MessagePack, TOML, CSV, binary formats, etc.
- Derive macros for convenience
- Supports `no_std` environments

# Serde Derive Example

```
#[derive(serde::Serialize, serde::Deserialize, Debug, PartialEq, Eq)]
enum Shape {
    Point(Point),
    Line { from: Point, to: Point },
}

#[derive(serde::Serialize, serde::Deserialize, Debug, PartialEq, Eq)]
struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let shape = Shape::Line {
        from: Point { x: 1, y: 2 },
        to: Point { x: 3, y: 4 },
    };

    let serialized = serde_json::to_string_pretty(&shape).unwrap();
    println!("{}", serialized);
    assert_eq!(serde_json::from_str::<Shape>(&serialized).unwrap(), shape);
}
```

Output:

```
{
  "Line": {
    "from": {
      "x": 1,
      "y": 2
    },
    "to": {
      "x": 3,
      "y": 4
    }
  }
}
```

# Serde Derive Example: YAML

```
#[derive(serde::Serialize, serde::Deserialize, Debug, PartialEq, Eq)]
enum Shape {
    Point(Point),
    Line { from: Point, to: Point },
}

#[derive(serde::Serialize, serde::Deserialize, Debug, PartialEq, Eq)]
struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let shape = Shape::Line {
        from: Point { x: 1, y: 2 },
        to: Point { x: 3, y: 4 },
    };

    let serialized = serde_yaml::to_string(&shape).unwrap();
    println!("{}", serialized);
    assert_eq!(serde_yaml::from_str::<Shape>(&serialized).unwrap(), shape);
}
```

YAML Output:

```
---
Line:
  from:
    x: 1
    y: 2
  to:
    x: 3
    y: 4
```

# Serde Derive: Attributes

Derived Serialize/Deserialize implementations can be customized through `#[serde]` attributes, e.g.:

- Rename containers or fields:

```
#[derive(serde::Serialize, serde::Deserialize)]
#[serde(rename = "GreatExample", rename_all = "camelCase")]
struct Example {
    field_one: u32,           // struct name will be changed to GreatExample
    some_other_field: u32,    // will be renamed to fieldOne
    #[serde(rename = "lastField")]
    field_three: bool,        // will be renamed to lastField
}
```

# Serde Derive: Attributes

Derived Serialize/Deserialize implementations can be customized through `#[serde]` attributes, e.g.:

- Rename containers or fields:

```
#[derive(serde::Serialize, serde::Deserialize)]
#[serde(rename = "GreatExample", rename_all = "camelCase")]
struct Example {
    field_one: u32,           // struct name will be changed to GreatExample
    some_other_field: u32,    // will be renamed to fieldOne
    #[serde(rename = "lastField")]
    field_three: bool,        // will be renamed to someOtherField
}
```

- Don't serialize None fields:

```
#[derive(serde::Serialize, serde::Deserialize)]
struct Example {
    #[serde(skip_serializing_if = "Option::is_none")]
    field: Option<u32>,
}
```

# Serde Derive: More Attributes

- Flatten nested structs:

```
#[derive(serde::Serialize, serde::Deserialize)]
struct Example {
    field_1: u32,
    #[serde(flatten)]
    field_2: Field2,
}

#[derive(serde::Serialize, serde::Deserialize)]
struct Field2 {
    field_a: bool,
    field_b: String,
}
```

With flatten:

```
{
  "field_1": 42,
  "field_a": true,
  "field_b": "Hello"
}
```

Without flatten:

```
{
  "field_1": 42,
  "field_2": {
    "field_a": true,
    "field_b": "Hello"
  }
}
```



# Serde Derive: More Attributes

- Flatten nested structs:

```
#[derive(serde::Serialize, serde::Deserialize)]
struct Example {
    field_1: u32,
    #[serde(flatten)]
    field_2: Field2,
}

#[derive(serde::Serialize, serde::Deserialize)]
struct Field2 {
    field_a: bool,
    field_b: String,
}
```

- Fill default values on deserialization:

```
#[derive(serde::Serialize, serde::Deserialize)]
struct Example {
    #[serde(default)]
    name: String,           // fills with empty string if not present
}
```

With flatten:

```
{
  "field_1": 42,
  "field_a": true,
  "field_b": "Hello"
}
```

Without flatten:

```
{
  "field_1": 42,
  "field_2": {
    "field_a": true,
    "field_b": "Hello"
  }
}
```

# The Serialize Trait

```
pub trait Serialize {  
    fn serialize<S>(&self, serializer: S) → Result<S::Ok, S::Error>  
        where S: Serializer;  
}
```

- Implemented for Rust types that can be serialized
  - maps Rust types to *serde data model*
- Serialize trait is independent of target format
  - generic parameter *S* allows passing in any compatible serializer
- Standard implementations can be derived
  - derive behavior can be customized through attributes
  - full manual implementation is possible too

# Serde Data Model

Most Rust types have corresponding types in the serde data model, e.g.:

- all primitive types, e.g. `i64`, floats, `bool`, `char`
- UTF8 strings, byte arrays, `Vec`, `HashMap`
- structs, enums, tuples

→ Serialization is often a trivial mapping

# Serde Data Model

Most Rust types have corresponding types in the serde data model, e.g.:

- all primitive types, e.g. `i64`, floats, `bool`, `char`
- UTF8 strings, byte arrays, `Vec`, `HashMap`
- structs, enums, tuples

→ Serialization is often a trivial mapping

---

Serialize implementations can **map to a different serde type** if desired. For example:

- `std::ffi::OsString` represents a platform-native string
  - not guaranteed to be UTF8-encoded, so serializing as serde string would be invalid
  - serialize as ASCII **byte array** on UNIX
  - serialize as **sequence of 16-bit UTF-16** values on Windows

# Manual Serialize implementation

```
struct Point {  
    x: i32,  
    y: i32,  
}  
  
impl serde::Serialize for Point {  
    fn serialize<S>(&self, serializer: S,) → Result<S::Ok, S::Error>  
    where  
        S: serde::Serializer,  
    {  
        let mut s = serializer.serialize_struct(  
            "Point",  
            2, // number of fields  
        )?;  
        s.serialize_field("x", &self.x)?;  
        s.serialize_field("y", &self.y)?;  
        s.end()  
    }  
}
```

# Example: Change type on Serialize

Serialize struct as string:

```
struct Point {  
    x: i32,  
    y: i32,  
}  
  
impl serde::Serialize for Point {  
    fn serialize<S>(&self, serializer: S) → Result<S::Ok, S::Error>  
    where  
        S: serde::Serializer,  
    {  
        let string_representation = format!("{}", self.x, self.y);  
        serializer.serialize_str(&string_representation)  
    }  
}
```

# The Deserialize Trait

```
pub trait Deserialize<'de>: Sized {  
    fn deserialize<D>(deserializer: D) → Result<Self, D::Error>  
        where D: Deserializer<'de>;  
}
```

- Goal: Drive the given Deserializer to map raw data into the serde data model
- Deserializer provides two kinds of entry points:
  - **deserialize\_any** for self-describing formats (e.g. JSON)
    - instructs the deserializer to read the data type from the raw data
    - panics for formats that are not self-describing (e.g. Cap'n Proto)
  - **typed deserialize\_\* functions**, e.g. `deserialize_f32` or `deserialize_string`

# The Deserialize Trait

```
pub trait Deserialize<'de>: Sized {  
    fn deserialize<D>(deserializer: D) → Result<Self, D::Error>  
        where D: Deserializer<'de>;  
}
```

- Goal: Drive the given Deserializer to map raw data into the serde data model
- Deserializer provides two kinds of entry points:
  - **deserialize\_any** for self-describing formats (e.g. JSON)
    - instructs the deserializer to read the data type from the raw data
    - panics for formats that are not self-describing (e.g. Cap'n Proto)
  - **typed deserialize\_\* functions**, e.g. `deserialize_f32` or `deserialize_string`
- Deserialize methods take a **Visitor** implementation that constructs the target value
  - translates the serde data format into a Rust type
  - gives additional type hints for nested data, e.g. structs



# Manual Deserialize Implementation

```
struct Number(u32);

impl<'de> serde::Deserialize<'de> for Number {
    fn deserialize<D>(deserializer: D) -> Result<Self, D::Error>
    where
        D: serde::Deserializer<'de>,
    {
        let visitor = NumberVisitor;
        deserializer.deserialize_u32(visitor)
    }
}

struct NumberVisitor;

impl<'de> Visitor<'de> for NumberVisitor {...}           // explained later
```

- we tell the deserializer that we expected a u32 value
  - allows deserializing from formats that are not self-describing

# Example: deserialize\_any

```
enum NumberOrString {  
    Number(u32),  
    String(String),  
}  
  
impl<'de> serde::Deserialize<'de> for NumberOrString {  
    fn deserialize<D>(deserializer: D) → Result<Self, D::Error>  
    where  
        D: serde::Deserializer<'de>,  
    {  
        let visitor = NumberOrStringVisitor;  
        deserializer.deserialize_any(visitor)  
    }  
}  
  
struct NumberOrStringVisitor;  
impl<'de> Visitor<'de> for NumberOrStringVisitor {...} // explained later
```

- only works with self-describing formats, e.g. JSON
- panics when used with other formats, e.g. Cap'n Proto

# The Visitor trait

```
pub trait Visitor<'de>: Sized {
    type Value;          // the target type that we want to deserialize to
    fn expecting(&self, formatter: &mut Formatter<'_>) -> Result;      // only required method

    fn visit_bool<E>(self, v: bool) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_i8<E>(self, v: i8) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_i16<E>(self, v: i16) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    ...
    fn visit_f64<E>(self, v: f64) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_char<E>(self, v: char) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_str<E>(self, v: &str) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_borrowed_str<E>(self, v: &'de str) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    fn visit_string<E>(self, v: String) -> Result<Self::Value, E> where E: serde::de::Error { ... }
    ...
    fn visit_seq<A>(self, seq: A) -> Result<Self::Value, A::Error>
        where A: SeqAccess<'de> { ... }
    fn visit_map<A>(self, map: A) -> Result<Self::Value, A::Error>
        where A: MapAccess<'de> { ... }
    fn visit_enum<A>(self, data: A) -> Result<Self::Value, A::Error>
        where A: EnumAccess<'de> { ... }
}
```

# Example: Creating a NumberVisitor

```
struct NumberVisitor;

impl<'de> serde::de::Visitor<'de> for NumberVisitor {
    type Value = Number;

    fn expecting(&self, formatter: &mut std::fmt::Formatter) → std::fmt::Result {
        formatter.write_str("an unsigned 32-bit integer")
    }

    fn visit_u32<E>(self, v: u32) → Result<Self::Value, E>
    where
        E: serde::de::Error,
    {
        Ok(Number(v))
    }
}
```

- all the `visit_*` methods are optional → deserialization error by default
- we expect a `u32`, so we only need to implement `visit_u32`, right? → **wrong!**

# Implementing `visit_u32` is not enough

The following code panics:

[Try it yourself on the playground](#)

```
let deserialized: Number = serde_json::from_str("42").unwrap();  
// → Error("invalid type: integer '42', expected an unsigned 32-bit integer", line: 1, column: 2)
```

# Implementing `visit_u32` is not enough

The following code panics:

[Try it yourself on the playground](#)

```
let deserialized: Number = serde_json::from_str("42").unwrap();  
// → Error("invalid type: integer '42', expected an unsigned 32-bit integer", line: 1, column: 2)
```

- the message after "expected" is from our `expected()` implementation
- the *"invalid type"* error message is coming from the default implementation of `visit_u64`
  - why is it called instead of `visit_u32`?
  - remember, we called `deserialize_u32` in our `Deserialize` implementation:

```
impl<'de> serde::Deserialize<'de> for Number {  
    fn deserialize<D>(deserializer: D) → Result<Self, D::Error> where D: serde::Deserializer<'de> {  
        let visitor = NumberVisitor;  
        deserializer.deserialize_u32(visitor)    // → leads to a visit_u64 call somehow  
    }  
}
```

# Fixing the NumberVisitor Example

- The `deserialize_*` methods are only a type hint
  - the `Deserializer` is allowed to call other `visit_*` methods
- There are no separate `u8`, `u16`, `u32`, and `u64` types in JSON
  - only a general unsigned type
  - → the `serde_json` `Deserializer` will always call `visit_u64` for all unsigned integers

# Fixing the NumberVisitor Example

- The `deserialize_*` methods are only a type hint
  - the `Deserializer` is allowed to call other `visit_*` methods
- There are no separate `u8`, `u16`, `u32`, and `u64` types in JSON
  - only a general unsigned type
  - → the `serde_json` `Deserializer` will always call `visit_u64` for all unsigned integers
- **Fix:** Add a `visit_u64` implementation to our `NumberVisitor`:


```
fn visit_u64<E>(self, v: u64) → Result<Self::Value, E> where E: serde::de::Error {  
    match u32::try_from(v) {  
        Ok(v) ⇒ Ok(Number(v)),  
        Err(_) ⇒ Err(E::custom("number is too large")),  
    }  
}
```

- Also: implement similar methods, e.g. `visit_u8`, `visit_i32`, and maybe `visit_f32`



# Dora ROS2 Bridge

- ROS2 is a mature framework for robot software development
- A bridge to ROS2 has many advantages for Dora
  - Use ROS2 tooling for Dora
    - For example, record and replay Dora messages using rosbag
  - Compatibility with existing ROS2 systems
    - Bridge allows partial migration to Dora
  - Drivers written for ROS2
    - Dora nodes can use them through bridge
- Use `ros2-client` crate to
  - communicate with ROS2 nodes via DDS
  - serialize and deserialize messages in CDR format



# CDR Format

- short for "Common Data Representation"
- binary format
- not self-describing
  - struct field names etc. are not stored
  - we need to know exact type for deserialization

# CDR Format

- short for "Common Data Representation"
- binary format
- not self-describing
  - struct field names etc. are not stored
  - we need to know exact type for deserialization
- ROS2 describes message format in .msg files
  - For example, the Vector3.msg file looks like this:

```
float64 x  
float64 y  
float64 z
```

- Format is relatively simple → we can parse it
- Base parser on existing `rc1rust-msg-gen` crate

# Serializing to dynamic target type

- When using Python nodes, we only know the available ROS2 types at runtime
  - because the `dora-daemon` and `dora-runtime` are precompiled on different machines
  - we still need to serialize them to the **correct ROS2 type**
- Approach:
  - parse ROS2 `.msg` files on initialization
  - supply ROS2 type name on topic creation function → look up type info from `.msg` file
  - on `publish`, use type info to serialize Python object to correct ROS2 type

# Serializing to dynamic target type

- When using Python nodes, we only know the available ROS2 types at runtime
  - because the dora-daemon and dora-runtime are precompiled on different machines
  - we still need to serialize them to the **correct ROS2 type**
- Approach:
  - parse ROS2 .msg files on initialization
  - supply ROS2 type name on topic creation function → look up type info from .msg file
  - on publish, use type info to serialize Python object to correct ROS2 type

## Example:

```
turtle_twist_topic = ros2_node.create_topic("/turtle1/cmd_vel", "geometry_msgs::Twist", topic_qos)
twist_writer = ros2_node.create_publisher(turtle_twist_topic)

# serializes 'value' object as 'geometry_msgs::Twist' message
twist_writer.publish(value)
```

# Dynamic Serialize Implementation

```
pub struct TypedValue<'a> {  
    pub value: &'a arrow::ArrayData,  
    pub type_info: &'a TypeInfo,  
}  
  
impl serde::Serialize for TypedValue<'_> {  
    fn serialize<S>(&self, serializer: S) → Result<S::Ok, S::Error>  
    where  
        S: serde::Serializer,  
    {  
        match &self.type_info.data_type {  
            DataType::UInt8 ⇒ {  
                let number = convert_to_u8(self.value);    // → this function could e.g. do integer tpe conversion  
                serializer.serialize_u8(number)  
            }  
            DataType::UInt16 ⇒ {...}  
            ...  
        }  
    }  
}
```

# Dynamic Serialize Implementation (2)

```
match &self.type_info.data_type {  
    // ...continued  
    DataType::List(item_info) => {  
        let list_array: ListArray = self.value.clone().into();  
        let mut s = serializer.serialize_seq(Some(list_array.values.len()))?;  
  
        for item_value in list_array.iter() {  
            // leads to recursive call  
            let element = TypedValue {  
                value: &item_value,  
                type_info: &TypeInfo {  
                    data_type: item_info.data_type().clone(),  
                    // ..omitted  
                },  
            };  
            s.serialize_element(&element)?;  
        }  
        s.end()  
    }  
}
```

# Dynamic deserialization

- We also want to deserialize ROS2 messages
  - CDR format is not self-describing → cannot use `deserialize_any`
  - we need to use correct `deserialize_*` for each deserialization step
  - → supply type name on subscribe call → get type info of expected type from `.msg` file



# Dynamic deserialization

- We also want to deserialize ROS2 messages
  - CDR format is not self-describing → cannot use `deserialize_any`
  - we need to use correct `deserialize_*` for each deserialization step
  - → supply type name on subscribe call → get type info of expected type from `.msg` file
- Challenge: No way to use runtime type info in `Deserialize` implementation:

```
impl<'de> serde::Deserialize<'de> for OwnedTypedValue {  
    fn deserialize<D>(deserializer: D) → Result<Self, D::Error>  
    where  
        D: serde::Deserializer<'de>,  
        {  
            match type_info.data_type { // → how to get access to type_info here?  
                DataType::UInt32 ⇒ deserializer.deserialize_u32(NumberVisitor),  
                DataType::List(item_info) ⇒ {...}  
                ...  
            }  
        }  
}
```

# Using DeserializeSeed

- It's not possible to use runtime data in Deserialize trait
- Fortunately, there is an alternative **DeserializeSeed** trait for this case:

```
pub trait DeserializeSeed<'de>: Sized {  
    type Value;  
  
    fn deserialize<D>(self, deserializer: D) → Result<Self::Value, D::Error>  
        where D: Deserializer<'de>;  
}
```

- runtime context can be stored in self object
- deserialize method returns Self::Value (instead of returning Self)

# DeserializeSeed Example

```
pub struct TypedDeserializer {
    type_info: TypeInfo,
}

impl<'de> serde::de::DeserializeSeed<'de> for TypedDeserializer {
    type Value = arrow::ArrayData;

    fn deserialize<D>(self, deserializer: D) → Result<Self::Value, D::Error>
    where
        D: serde::Deserializer<'de>,
    {
        match self.type_info.data_type {
            DataType::UInt32 ⇒ deserializer.deserialize_u32(NumberVisitor),
            DataType::List(item_info) ⇒ {...}
            ...
        }
    }
}
```

# Summary

- The `serde` crate provides format-independent **Serialize** and **Deserialize** traits
  - Also: **DeserializeSeed** for deserialization with additional runtime state
  - Traits map between `serde` data model and Rust types
  - External crates for mapping `serde` data model to **different serialization formats**
- **Derive macros** make it easy to add `serde` support for custom structs/enums
  - Attributes allow modifying serialization format, e.g. renames, skipping fields, flattening structs
- Manual `Deserialize` implementations require custom **Visitor**
  - **deserialize\_any** only works with self-describing data formats
  - no guarantee that deserializer follows `deserialize_*` type hint
- Dora supports reading and writing **typed ROS2 messages** using custom `serde` implementations