

DORA-RS

On Async Runtime

github.com/dora-rs/dora

Xavier Tao

TOKIO RUNTIME

Very powerful async engine that is able to run green threads, running highly-concurrent tasks!

```
#[tokio::main]
async fn main() -> eyre::Result<()> {
```

ASYNC DESIGN: EVENT STREAM

```
enum Event {  
    Input(InputEvent),  
    InputsStopped,  
    Operator {  
        id: OperatorId,  
        event: OperatorEvent,  
    },  
}
```

ASYNCHRONOUS DESIGN: EVENT STREAM

```
enum Event {  
  Input(InputEvent),  
  InputsStopped,  
  Operator {  
    id: OperatorId,  
    event: OperatorEvent,  
  },  
}
```

```
let inputs = subscribe(&topics, communication).await?;  
let input_events = inputs.map(Event::Input)  
  .chain(stream::once(async { Event::InputsStopped }));
```

ASYNC DESIGN: EVENT STREAM

```
enum Event {  
    Input(InputEvent),  
    InputsStopped,  
    Operator {  
        id: OperatorId,  
        event: OperatorEvent,  
    },  
}
```

```
let inputs = subscribe(&topics, communication).await?;  
let input_events = inputs.map(Event::Input)  
    .chain(stream::once(async { Event::InputsStopped }));
```

```
let mut operator_events = tokio_stream::StreamMap::new();  
for operator in operators {  
    operator_events.insert(operator.id, spawn_operator().await);  
}  
let op_events = operator_events.map(|(id, event)| Event::Operator { id, event });
```

ASYNC DESIGN: EVENT STREAM

```
enum Event {  
    Input(InputEvent),  
    InputsStopped,  
    Operator {  
        id: OperatorId,  
        event: OperatorEvent,  
    },  
}
```

```
let inputs = subscribe(&topics, communication).await?;  
let input_events = inputs.map(Event::Input)  
    .chain(stream::once(async { Event::InputsStopped }));
```

```
let mut operator_events = tokio_stream::StreamMap::new();  
for operator in operators {  
    operator_events.insert(operator.id, spawn_operator().await);  
}  
let op_events = operator_events.map(|(id, event)| Event::Operator { id, event });
```

```
use futures_concurrency::Merge;  
let mut events = (input_events, op_events).merge();  
  
// now we have a normal while loop with a match instead of needing `select`  
while let Some(event) = events.next().await {  
    match event { ... }  
}
```

CHALLENGES WITH ASYNC RUST

- CPU-bound operations in async functions
- Cancellation of spawned tasks
- Error/panic propagation

CPU-BOUND OPERATIONS IN ASYNC FUNCTIONS


*To give a sense of scale of how much time is too much, a good rule of thumb is **no more than 10 to 100 microseconds** between each `.await`.*

	CPU-bound computation	Synchronous IO	Running forever
<code>spawn_blocking</code>	Suboptimal	OK	No
<code>rayon</code>	OK	No	No
Dedicated thread	OK	OK	OK

— Alice Ryhl, in *“[Async: What is blocking?](#)”*

SPAWNING CPU-BOUND TASKS

- **tokio::spawn_blocking**: Optimal for reducing tail latency
 - Pros
 - Directly available threadpool
 - Work well enough on most cases
 - Cons
 - Need adequate parametrisation (workers, threads, timeout...)
- **rayon**: Optimal for maximizing cpu usage
 - Pros
 - Do not need Arc for ∞ read
 - Stack allocation of tasks instead of heap
 - Recursive parallelisation
 - Cons
 - Add a layer of technology



CANCELLATION OF SPAWNED TASKS

- Dropping a `JoinHandle` detaches the task
→ keeps running in the background

```
let task = tokio::spawn(async {
    let mut items = Vec::new();
    for _ in 0..n {
        items.push(read_large_file().await);
    }
    let _ = result_tx.send(items).await;
});
do_something_else()?; // ~ `task` keeps running on error
let items = task.await;
```

CANCELLATION OF SPAWNED TASKS

- Dropping a `JoinHandle` detaches the task
→ keeps running in the background

```
let task = tokio::spawn(async {  
    let mut items = Vec::new();  
    for _ in 0..n {  
        items.push(read_large_file().await);  
    }  
    let _ = result_tx.send(items).await;  
});  
do_something_else()?; // ^ `task` keeps running on error  
let items = task.await;
```

- Workarounds
 - manual checks (e.g. is `result_tx` closed?)
→ error-prone and less efficient
 - use `smol` instead of `tokio` or `async_std`
 - use `FutureExt::remote_handle` everywhere

```
let (task, handle) = task.remote_handle();  
tokio::spawn(task);  
do_something_else()?; // dropping `handle` cancels task  
handle.await
```

ERROR/PANIC PROPAGATION

- Easy to accidentally discard a panic/error

```
tokio::spawn(async { panic!("foo") });  
tokio::spawn(async { Result::<(), u32>::Err(1) });
```

- neither a compiler warning nor a runtime error occurs (just some `stderr` output for the panic)
- same with `async_std`
- `must_use` warning with `smol` (tasks are canceled on drop)

ASYNC ANNOYANCES

- Compiler cannot infer error type in async blocks

```
let task = async {  
  value.context("I/O error"?;  
  Ok())  
  // ^^ cannot infer type for type parameter `E`  
  // workaround: `Result::<_, anyhow::Error>::Ok()`  
};
```

ASYNC ANNOYANCES

- Compiler cannot infer error type in async blocks

```
let task = async {  
  value.context("I/O error"?);  
  Ok()  
  // ^^ cannot infer type for type parameter `E`  
  // workaround: `Result::<_, anyhow::Error>::Ok()`  
};
```

- Clone boilerplate with `async move {}`

```
let id_clone = id.clone();  
let tx_clone = tx.clone();  
// etc ...  
let task = async move {  
  tx_clone.send(id_clone).await  
};  
// (we still need `id`, `tx`, etc. here)
```

(also applies to move closures)

DON'T USE `select` MACRO

- requires lots of pinning and fusing
- its custom syntax breaks tools such as `rust-analyzer` or `rustfmt`
- `FusedFuture` bound does not prevent poll after `Poll::Ready`

```
loop {  
    let mut fut = (&mut fut).fuse(); // ~ (should be outside of loop)  
    select! {  
        _ = fut => println!("foo"),  
    };  
}
```


DON'T USE `select` MACRO

- requires lots of pinning and fusing
- its custom syntax breaks tools such as `rust-analyzer` or `rustfmt`
- `FusedFuture` bound does not prevent poll after `Poll::Ready`

```
loop {  
    let mut fut = (&mut fut).fuse(); // ~ (should be outside of loop)  
    select! {  
        _ = fut => println!("foo"),  
    };  
}
```

- Use an event stream instead
 - as proposed by Yoshua Wuyts in *[“Futures Concurrency III”](#)*

Q&A